




● NVIDIA



H100 / H200
Hopper GPU (SXM5)



Blackwell B200
Dual-Die GPU




GB200 NVL72
Grace Blackwell

● AMD Instinct




MI300X
CDNA 3 Multi-Chiplet



MI325X / MI350
CDNA 4 AI Accelerator

● Cerebras Systems



WSE-2 / WSE-3
Wafer Scale Engine (4T Transistors)



CS-3 System
Wafer Scale Cluster Block

● Google Cloud TPU



TPU v5p
Pod Training ASIC




TPU v6 Trillium
6th Gen Matrix NPU

● AWS Annapurna



Trainium2
Scale-Out Training



Inferentia2
Low-Latency Inference

● Intel Habana




Gaudi 2
Deep Learning Matrix



Gaudi 3
Heterogeneous AI NPU

● Meta



MTIA v2
Artemis Inference

● Graphcore




Bow IPU Mk2
3D Wafer-on-Wafer

● Groq



GroqCard LPU
Deterministic SRAM Stream

● Etched



Sohu
Hardwired Transformer ASIC

● Tenstorrent



Wormhole
Tenix Spatial Mesh

● SambaNova



SN40L
Reconfigurable Dataflow

● Taalas & MatX



Taalas-1
Direct Silicon Logic



MatX MXC
Large-Core LLM Engine