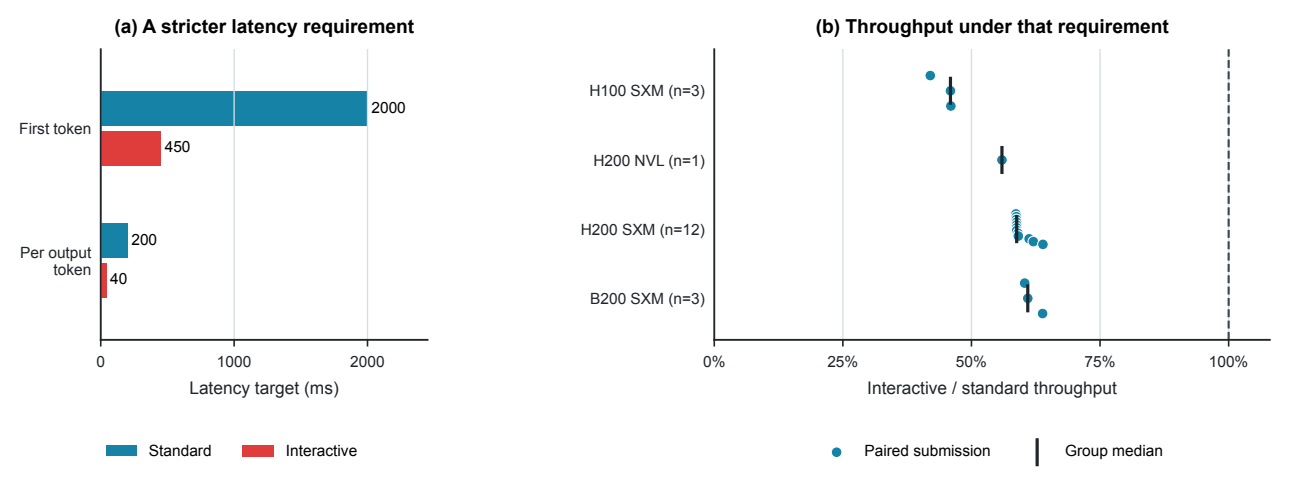
(a) A stricter latency requirement
First token
Per output token
2000
450
200
40
0
1000
2000
Latency target (ms)
Standard
Interactive
(b) Throughput under that requirement
H100 SXM (n=3)
H200 NVL (n=1)
H200 SXM (n=12)
B200 SXM (n=3)
0%
25%
50%
75%
100%
Interactive / standard throughput
Paired submission
Group median