



(a) Zero-Point Baseline Distortion in Agentic Benchmarks  
(Empirical Audit Data from Zhu et al. 2025 [arXiv:2507.02825])



(b) Hardware RTL Synthesis: Contamination vs. Accuracy  
(Empirical Evaluation from Wang et al. 2025 [arXiv:2503.13572])



Empirical Benchmark Integrity Audit: (a) Null agents (Do-Nothing, DB-Dump, Zip-Overwrite) score up to 38-100% on leading agentic suites, proving reported progress is distorted by uncalibrated zero-point floors and reward leaks (Zhu et al. 2025). (b) In Verilog RTL generation, high pass@1 rates strongly correlate with pretraining data contamination (CDD  $\geq 75\%$  in GPT-4o/3.5), while clean models achieve  $<10\%$  (Wang et al. 2025). Rigorous hardware evaluation mandates ephemeral sandbox isolation, trivial-baseline subtraction, and sealed multi-corner signoff.