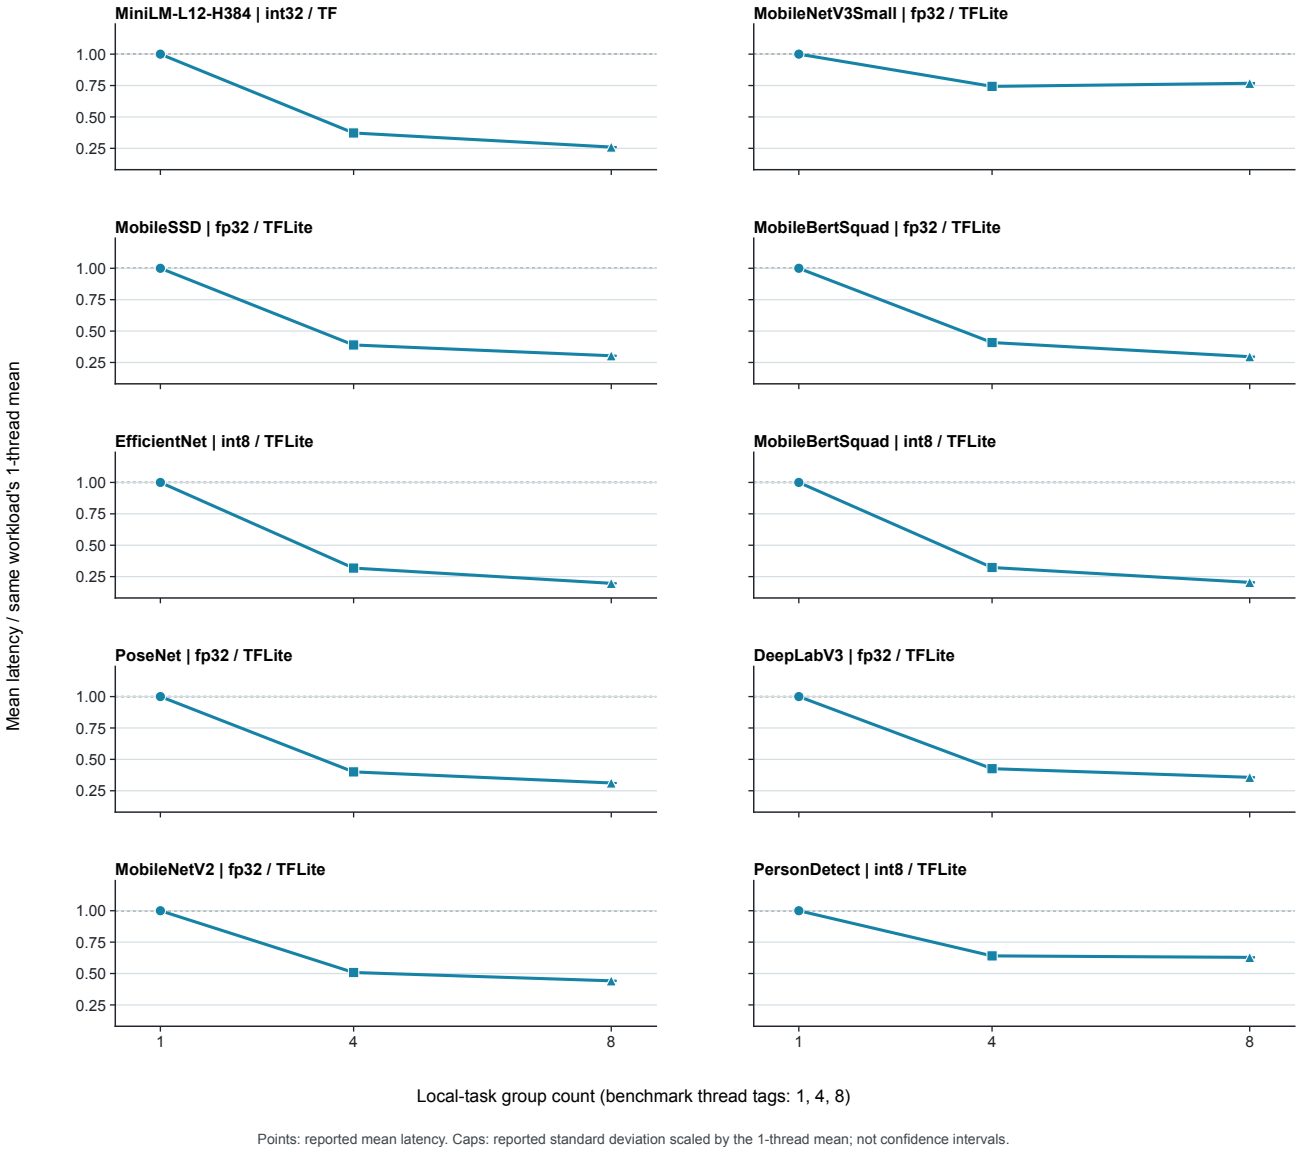
Thread scaling depends on the workload
IREE bot report 1559 · GCP c2-standard-16 label · Cascade Lake · LLVM-CPU local-task
MiniLM-L12-H384 | int32 / TF
MobileNetV3Small | fp32 / TFLite
MobileSSD | fp32 / TFLite
MobileBertSquad | fp32 / TFLite
EfficientNet | int8 / TFLite
MobileBertSquad | int8 / TFLite
PoseNet | fp32 / TFLite
DeepLabV3 | fp32 / TFLite
MobileNetV2 | fp32 / TFLite
PersonDetect | int8 / TFLite
1.00
0.75
0.50
0.25
1
4
8
Mean latency / same workload's 1-thread mean
Local-task group count (benchmark thread tags: 1, 4, 8)
Points: reported mean latency. Caps: reported standard deviation scaled by the 1-thread mean; not confidence intervals.